

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
```

```
df = pd.read_csv("Housing.csv")
display(df)
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking
0	13300000	7420	4	2	3	yes	no	no	no	yes	2
1	12250000	8960	4	4	4	yes	no	no	no	yes	3
2	12250000	9960	3	2	2	yes	no	yes	no	no	2
3	12215000	7500	4	2	2	yes	no	yes	no	yes	3
4	11410000	7420	4	1	2	yes	yes	yes	no	yes	2
...
540	1820000	3000	2	1	1	yes	no	yes	no	no	2
541	1767150	2400	3	1	1	no	no	no	no	no	0
542	1750000	3620	2	1	1	yes	no	no	no	no	0
543	1750000	2910	3	1	1	no	no	no	no	no	0
544	1750000	3850	3	1	2	yes	no	no	no	no	0

545 rows × 13 columns



Probleme 1a

```
#accessing columns that were asked
#loc = index by name and name is string
reduced_df = df.loc[:,["area", "bedrooms", "bathrooms", "stories", "parking"]]
price = df.iloc[:,0] #iloc = index by number
```

```
maxes = []
# for i in range(5):
#     maxes.append(max(reduced_df.iloc[:,i]))
#     reduced_df.iloc[:,i] = reduced_df.iloc[:,i]/maxes[-1]

# maxes.append(max(price))
# price = price/maxes[-1]
# reduced_df.iloc[:,0] = reduced_df.iloc[:,0]/max(reduced_df.iloc[:,0])
print(price)
print(reduced_df.iloc[:,:])
print(maxes)
```

```
0      13300000
1      12250000
2      12250000
3      12215000
4      11410000
...
540     1820000
541     1767150
542     1750000
543     1750000
544     1750000
Name: price, Length: 545, dtype: int64
   area  bedrooms  bathrooms  stories  parking
0   7420         4          2         3         2
1   8960         4          4         4         3
2   9960         3          2         2         2
3   7500         4          2         2         3
4   7420         4          1         2         2
..  ...      ...      ...      ...      ...
540  3000         2          1         1         2
541  2400         3          1         1         0
```

542	3620	2	1	1	0
543	2910	3	1	1	0
544	3850	3	1	2	0

```
[545 rows x 5 columns]
[]
```

```
X_train, X_test, y_train, y_test = train_test_split(reduced_df.values, price.values, test_size = 0.2, random_state =42)
```

```
print('features:\n', X_train[0]) #shows all 5 X values
print('example price:', y_train[0])
```

```
features:
[6000  3  2  4  1]
example price: 7525000
```

```
# 5x5
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
m = len(X_train)
X_train = X_train.reshape(m,5)
print('Xvals: \n', X_train[:5,:])
m
```

```
Xvals:
[[6000  3  2  4  1]
 [7200  3  2  1  3]
 [3816  2  1  1  2]
 [2610  3  1  2  0]
 [3750  3  1  2  0]]
436
```

```
m = len(X_train)
n = len(X_test)
X_train_1 = np.ones((m,1)) #1 column
X_test_1 = np.ones((n,1)) #1 column
X_train_1[:1], len(X_train_1) #1 row
```

```
(array([[1.]]), 436)
```

```
X_train = np.hstack((X_train_1, X_train))
X_test = np.hstack((X_test_1, X_test))
X_train[:5]
```

```
array([[1.000e+00, 6.000e+03, 3.000e+00, 2.000e+00, 4.000e+00, 1.000e+00],
       [1.000e+00, 7.200e+03, 3.000e+00, 2.000e+00, 1.000e+00, 3.000e+00],
       [1.000e+00, 3.816e+03, 2.000e+00, 1.000e+00, 1.000e+00, 2.000e+00],
       [1.000e+00, 2.610e+03, 3.000e+00, 1.000e+00, 2.000e+00, 0.000e+00],
       [1.000e+00, 3.750e+03, 3.000e+00, 1.000e+00, 2.000e+00, 0.000e+00]])
```

```
theta = np.zeros(6)
iterations = 1500
alpha = .000000001
theta.shape
```

```
(6,)
```

```
def compute_loss(X, Y, theta):
    predictions = X.dot(theta)
    errors = np.subtract(predictions, Y)
    sqrErrors = np.square(errors)
    J = 1/(2*m) * np.sum(sqrErrors)
    return J
```

```
cost = compute_loss(X_train,y_train, theta)
print('The cost for given values of theta_0 and theta_1 = ', cost)
```

```
The cost for given values of theta_0 and theta_1 = 12617396203243.809
```

```
def gradient_descent(X, Y, theta, alpha, iterations, train=True):
    cost_history = np.zeros(iterations)

    for i in range(iterations):
        predictions = X.dot(theta)
```

```

errors = np.subtract(predictions, Y)
if train:
    sum_delta = (alpha/m)*X.transpose().dot(errors)
    theta = theta - sum_delta
    cost_history[i] = compute_loss(X, Y, theta)

return theta, cost_history

theta, cost_history = gradient_descent(X_train, y_train, theta, alpha, iterations)
theta, test_history = gradient_descent(X_test, y_test, theta, alpha, iterations, train = False)

print('Final value of theta =', theta)
print('cost_history = ', cost_history[2:])
len(cost_history)

```

```


Final value of theta = [1.94414246e-01 8.30456226e+02 6.23040970e-01 2.92262578e-01
4.24703458e-01 1.55536615e-01]
cost_history = [1.24111673e+13 1.23432850e+13 1.22758284e+13 ... 1.59026477e+12
1.59025920e+12 1.59025365e+12]
1500

```

```

plt.plot(cost_history, color = 'pink', label = 'training data')
plt.plot(test_history, color = 'blue', label = 'testing data')

```



Problem 1b

```

#accessing columns that were asked
#loc = index by name and name is string
reduced_df = df.loc[:,["area", "bedrooms", "bathrooms", "stories", "mainroad", "guestroom", "basement", "hotwaterheating", "airconditioning",
price = df.iloc[:,0] #iloc = index by number

# arr = [1, 2, 3]
# print([i for i in arr])
# print([i**2 for i in arr])
# print([i for i in arr if i%2 == 1])

for column_index in [4, 5, 6, 7, 8, 10]:
    reduced_df.iloc[:,column_index] = reduced_df.iloc[:,column_index].map({"yes": 1, "no": 0})
display(reduced_df.head())

maxes = []

# for i in range(11):
#     maxes.append(max(reduced_df.iloc[:,i]))
#     reduced_df.iloc[:,i] = reduced_df.iloc[:,i]/maxes[-1]

# maxes.append(max(price))
# price = price/maxes[-1]
# reduced_df.iloc[:,0] = reduced_df.iloc[:,0]/max(reduced_df.iloc[:,0])
print(price)
print(reduced_df.iloc[:,:])
print(maxes)

```

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea
0	7420	4	2	3	1	0	0	0	1	2	1
1	8960	4	4	4	1	0	0	0	1	3	0
2	9960	3	2	2	1	0	1	0	0	2	1
3	7500	4	2	2	1	0	1	0	1	3	1
4	7420	4	1	2	1	1	1	0	1	2	0



```
0 13300000
1 12250000
2 12250000
3 12215000
4 11410000
```

```
...
540 1820000
541 1767150
542 1750000
543 1750000
544 1750000
```

Name: price, Length: 545, dtype: int64

```
area bedrooms bathrooms stories mainroad guestroom basement \
0 7420 4 2 3 1 0 0
1 8960 4 4 4 1 0 0
2 9960 3 2 2 1 0 1
3 7500 4 2 2 1 0 1
4 7420 4 1 2 1 1 1
.. ...
540 3000 2 1 1 1 0 1
541 2400 3 1 1 0 0 0
542 3620 2 1 1 1 0 0
543 2910 3 1 1 0 0 0
544 3850 3 1 2 1 0 0
```

```
hotwaterheating airconditioning parking prefarea
0 0 1 2 1
1 0 1 3 0
2 0 0 2 1
3 0 1 3 1
4 0 1 2 0
```

```
X_train, X_test, y_train, y_test = train_test_split(reduced_df.values, price.values, test_size = 0.2, random_state =42)
```

```
544 0 0 0 0 0
```

```
print('features:\n', X_train[0]) #shows all 5 X values
print('example price:', y_train[0])
```

```
features:
[6000 3 2 4 1 0 0 0 1 1 0]
example price: 7525000
```

```
# 5x5
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
m = len(X_train)
X_train = X_train.reshape(m,11)
print('Xvals: \n', X_train[:11,:])
m
```

```
Xvals:
[[6000 3 2 4 1 0 0 0 1 1 0]
[7200 3 2 1 1 0 1 0 1 3 0]
[3816 2 1 1 1 0 1 0 1 2 0]
[2610 3 1 2 1 0 1 0 0 0 1]
[3750 3 1 2 1 0 0 0 0 0 0]
[5010 3 1 2 1 0 1 0 0 0 0]
[3850 3 1 2 1 0 0 0 0 0 0]
[5076 3 1 1 0 0 0 0 0 0 0]
[3180 3 1 1 0 0 0 0 0 0 0]
[6020 3 1 1 1 0 0 0 0 0 0]
[3120 3 1 2 0 0 1 1 0 0 0]]
```

436

```
m = len(X_train)
n = len(X_test)
X_train_1 = np.ones((m,1)) #1 column
X_test_1 = np.ones((n,1)) #1 column
X_train_1[:1], len(X_train_1) #1 row
```

```

(array([[1.]]), 436)

X_train = np.hstack((X_train_1, X_train))
X_test = np.hstack((X_test_1, X_test))
X_train[:5]

array([[1.000e+00, 6.000e+03, 3.000e+00, 2.000e+00, 4.000e+00, 1.000e+00,
        0.000e+00, 0.000e+00, 0.000e+00, 1.000e+00, 1.000e+00, 0.000e+00],
       [1.000e+00, 7.200e+03, 3.000e+00, 2.000e+00, 1.000e+00, 1.000e+00,
        0.000e+00, 1.000e+00, 0.000e+00, 1.000e+00, 3.000e+00, 0.000e+00],
       [1.000e+00, 3.816e+03, 2.000e+00, 1.000e+00, 1.000e+00, 1.000e+00,
        0.000e+00, 1.000e+00, 0.000e+00, 1.000e+00, 2.000e+00, 0.000e+00],
       [1.000e+00, 2.610e+03, 3.000e+00, 1.000e+00, 2.000e+00, 1.000e+00,
        0.000e+00, 1.000e+00, 0.000e+00, 0.000e+00, 0.000e+00, 1.000e+00],
       [1.000e+00, 3.750e+03, 3.000e+00, 1.000e+00, 2.000e+00, 1.000e+00,
        0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00, 0.000e+00]])

theta = np.zeros(12)
iterations = 1500
alpha = .000000001
theta.shape

(12,)

def compute_loss(X, Y, theta):
    predictions = X.dot(theta)
    errors = np.subtract(predictions, Y)
    sqrErrors = np.square(errors)
    J = 1/(2*m) * np.sum(sqrErrors)
    return J

cost = compute_loss(X_train,y_train, theta)
print('The cost for given values of theta_0 and theta_1 = ', cost)

The cost for given values of theta_0 and theta_1 = 12617396203243.809

def gradient_descent(X, Y, theta, alpha, iterations, train = True):
    cost_history = np.zeros(iterations)

    for i in range(iterations):
        predictions = X.dot(theta)
        errors = np.subtract(predictions, Y)
        sum_delta = (alpha/m)*X.transpose().dot(errors)
        theta = theta - sum_delta
        cost_history[i] = compute_loss(X, Y, theta)

    return theta, cost_history

theta, cost_history = gradient_descent(X_train, y_train,theta, alpha, iterations)
print('Final value of theta =', theta)
print('cost_history = ', cost_history[2:])
len(cost_history)

Final value of theta = [1.94414238e-01 8.30456190e+02 6.23040945e-01 2.92262567e-01
 4.24703443e-01 1.72499519e-01 4.67442179e-02 9.14051009e-02
 1.43612303e-02 9.52339961e-02 1.55536609e-01 5.97006390e-02]
cost_history = [1.24111673e+13 1.23432850e+13 1.22758284e+13 ... 1.59026464e+12
 1.59025907e+12 1.59025353e+12]
1500

plt.plot(cost_history, color = 'pink', label = 'training data')
plt.plot(test_history, color = 'blue', label = 'testing data')

```

```
[<matplotlib.lines.Line2D at 0x7f3770565ca0>]
```



Problem 2a

```
def normalize_dataframe(df):
    numerical_cols = df.select_dtypes(include=[np.number]).columns
    df_normalized = df.copy()
    for col in numerical_cols:
        col_mean = df[col].mean()
        col_std = df[col].std()
        df_normalized[col] = (df[col] - col_mean) / col_std
    return df_normalized

#accessing columns that were asked
#loc = index by name and name is string
reduced_df = df.loc[:,["area", "bedrooms", "bathrooms", "stories", "parking"]]
price = df.iloc[:,0] #iloc = index by number

maxes = []
reduced_df = normalize_dataframe(reduced_df)
for i in range(5):
    maxes.append(max(reduced_df.iloc[:,i]))
    reduced_df.iloc[:,i] = reduced_df.iloc[:,i]/maxes[-1]

maxes.append(max(price))
price = price/maxes[-1]
reduced_df.iloc[:,0] = reduced_df.iloc[:,0]/max(reduced_df.iloc[:,0])
print(price)
print(reduced_df.iloc[:,:])
print(maxes)

0      1.000000
1      0.921053
2      0.921053
3      0.918421
4      0.857895
...
540    0.136842
541    0.132868
542    0.131579
543    0.131579
544    0.131579
Name: price, Length: 545, dtype: float64
   area  bedrooms  bathrooms  stories  parking
0  0.205391  0.340992   0.263016  0.544314  0.566428
1  0.344764  0.340992   1.000000  1.000000  1.000000
2  0.435266  0.011487   0.263016  0.088629  0.566428
3  0.212631  0.340992   0.263016  0.088629  1.000000
4  0.205391  0.340992  -0.105477  0.088629  0.566428
..  ...      ...      ...      ...      ...
540 -0.194629 -0.318017 -0.105477 -0.367057  0.566428
541 -0.248930  0.011487 -0.105477 -0.367057 -0.300716
542 -0.138517 -0.318017 -0.105477 -0.367057 -0.300716
543 -0.202774  0.011487 -0.105477 -0.367057 -0.300716
544 -0.117702  0.011487 -0.105477  0.088629 -0.300716

[545 rows x 5 columns]
[5.091585570242134, 4.111923842176972, 5.400846899372526, 2.529699688044935, 2.6769500506662025, 13300000]
```

```
X_train, X_test, y_train, y_test = train_test_split(reduced_df.values, price.values, test_size = 0.2, random_state =42)

print('features:\n', X_train[0]) #shows all 5 X values
print('example price:', y_train[0])

features:
[0.07687786 0.0114873 0.26301555 1.          0.13285601]
example price: 0.5657894736842105

# 5x5
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
m = len(X_train)
X_train = X_train.reshape(m,5)
```

```

print('Xvals: \n', X_train[:5,:])
m
Xvals:
[[ 0.07687786  0.0114873  0.26301555  1.          0.13285601]
 [ 0.18548046  0.0114873  0.26301555 -0.36705686  1.          ]
 [-0.12077888 -0.31801693 -0.10547667 -0.36705686  0.566428   ]
 [-0.2299245   0.0114873  -0.10547667  0.08862876 -0.30071599]
 [-0.12675203  0.0114873  -0.10547667  0.08862876 -0.30071599]]
436

m = len(X_train)
n = len(X_test)
X_train_1 = np.ones((m,1)) #1 column
X_test_1 = np.ones((n,1)) #1 column
X_train_1[:1], len(X_train_1) #1 row

(array([[1.]]), 436)

X_train = np.hstack((X_train_1, X_train))
X_test = np.hstack((X_test_1, X_test))
X_train[:5]

array([[ 1.          ,  0.07687786,  0.0114873 ,  0.26301555,  1.          ,
         0.13285601],
 [ 1.          ,  0.18548046,  0.0114873 ,  0.26301555, -0.36705686,
         1.          ],
 [ 1.          , -0.12077888, -0.31801693, -0.10547667, -0.36705686,
         0.566428   ],
 [ 1.          , -0.2299245 ,  0.0114873 , -0.10547667,  0.08862876,
        -0.30071599],
 [ 1.          , -0.12675203,  0.0114873 , -0.10547667,  0.08862876,
        -0.30071599]])

theta = np.zeros(6)
iterations = 1500
alpha = .01
theta.shape

(6,)

def compute_loss(X, Y, theta):
    predictions = X.dot(theta)
    errors = np.subtract(predictions, Y)
    sqrErrors = np.square(errors)
    J = 1/(2*m) * np.sum(sqrErrors)
    return J

cost = compute_loss(X_train,y_train, theta)
print('The cost for given values of theta_0 and theta_1 = ', cost)

The cost for given values of theta_0 and theta_1 =  0.07132905310217541

def gradient_descent(X, Y, theta, alpha, iterations, train=True):
    cost_history = np.zeros(iterations)

    for i in range(iterations):
        predictions = X.dot(theta)
        errors = np.subtract(predictions, Y)
        if train:
            sum_delta = (alpha/m)*X.transpose().dot(errors)
            theta = theta - sum_delta
        cost_history[i] = compute_loss(X, Y, theta)

    return theta, cost_history

theta, cost_history = gradient_descent(X_train, y_train,theta, alpha, iterations)
theta, test_history = gradient_descent(X_test, y_test,theta, alpha, iterations, train = False)


print('Final value of theta =', theta)
print('cost_history = ', cost_history[2:])
len(cost_history)

Final value of theta = [0.35598472 0.12594212 0.06378406 0.10164363 0.09412014 0.08882613]
cost_history = [0.06763439 0.0664511 0.06529114 ... 0.00441907 0.00441863 0.00441819]

```

1500

```
plt.plot(cost_history, color = 'pink', label = 'training data')
plt.plot(test_history, color = 'blue', label = 'testing data')
```



Problem 2b

```
#accessing columns that were asked
#loc = index by name and name is string
reduced_df = df.loc[:,["area", "bedrooms", "bathrooms", "stories", "mainroad", "guestroom", "basement", "hotwaterheating", "airconditioning",
price = df.iloc[:,0] #iloc = index by number

# arr = [1, 2, 3]
# print([i for i in arr])
# print([i**2 for i in arr])
# print([i for i in arr if i%2 == 1])

for column_index in [4, 5, 6, 7, 8, 10]:
    reduced_df.iloc[:,column_index] = reduced_df.iloc[:,column_index].map({"yes": 1, "no": 0})
display(reduced_df.head())

maxes = []

reduced_df = normalize_dataframe(reduced_df)

for i in range(11):
    maxes.append(max(reduced_df.iloc[:,i]))
    reduced_df.iloc[:,i] = reduced_df.iloc[:,i]/maxes[-1]

maxes.append(max(price))
price = price/maxes[-1]
reduced_df.iloc[:,0] = reduced_df.iloc[:,0]/max(reduced_df.iloc[:,0])
print(price)
print(reduced_df.iloc[:,:])
print(maxes)
```


	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea
0	7420	4	2	3	1	0	0	0	1	2	1
1	8960	4	4	4	1	0	0	0	1	3	0
2	9960	3	2	2	1	0	1	0	0	2	1
3	7500	4	2	2	1	0	1	0	1	3	1
4	7420	4	1	2	1	1	1	0	1	2	0



```
0 1.000000
1 0.921053
2 0.921053
3 0.918421
4 0.857895
```

```
...
540 0.136842
541 0.132868
542 0.131579
543 0.131579
544 0.131579
```

Name: price, Length: 545, dtype: float64

```
area bedrooms bathrooms stories mainroad guestroom basement \
0 0.205391 0.340992 0.263016 0.544314 1.000000 -0.216518 -0.539548
1 0.344764 0.340992 1.000000 1.000000 1.000000 -0.216518 -0.539548
```

```
X_train, X_test, y_train, y_test = train_test_split(reduced_df.values, price.values, test_size = 0.2, random_state =42)
```

```
4 0.205391 0.340992 -0.105477 0.088629 1.000000 1.000000 1.000000
```

```
print('features:\n', X_train[0]) #shows all 5 X values
```

```
print('example price:', y_train[0])
```

```
features:
[ 0.07687786  0.0114873  0.26301555  1.          1.          -0.21651786
 -0.53954802 -0.04807692  1.          0.13285601 -0.30695444]
example price: 0.5657894736842105
0 -0.048077  1.000000 0.566428 1.000000
```

```
# 5x5
```

```
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
```

```
m = len(X_train)
```

```
X_train = X_train.reshape(m,11)
```

```
print('Xvals: \n', X_train[:11,:])
```

```
m
```

```
Xvals:
[[ 0.07687786  0.0114873  0.26301555  1.          1.          -0.21651786
 -0.53954802 -0.04807692  1.          0.13285601 -0.30695444]
 [ 0.18548046  0.0114873  0.26301555 -0.36705686  1.          -0.21651786
  1.          -0.04807692  1.          1.          -0.30695444]
 [-0.12077888 -0.31801693 -0.10547667 -0.36705686  1.          -0.21651786
  1.          -0.04807692  1.          0.566428  -0.30695444]
 [-0.2299245  0.0114873 -0.10547667  0.08862876  1.          -0.21651786
  1.          -0.04807692 -0.46112601 -0.30071599  1.          ]
 [-0.12675203  0.0114873 -0.10547667  0.08862876  1.          -0.21651786
 -0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
 [-0.01271929  0.0114873 -0.10547667  0.08862876  1.          -0.21651786
  1.          -0.04807692 -0.46112601 -0.30071599 -0.30695444]
 [-0.11770181  0.0114873 -0.10547667  0.08862876  1.          -0.21651786
 -0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
 [-0.00674615  0.0114873 -0.10547667 -0.36705686 -6.07792208 -0.21651786
 -0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
 [-0.17833826  0.0114873 -0.10547667 -0.36705686 -6.07792208 -0.21651786
 -0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
 [ 0.0786879  0.0114873 -0.10547667 -0.36705686  1.          -0.21651786
 -0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
 [-0.18376839  0.0114873 -0.10547667  0.08862876 -6.07792208 -0.21651786
  1.          1.          -0.46112601 -0.30071599 -0.30695444]]
```

```
436
```

```
m = len(X_train)
```

```
n = len(X_test)
```

```
X_train_1 = np.ones((m,1)) #1 column
```

```
X_test_1 = np.ones((n,1)) #1 column
```

```
X_train_1[:5], len(X_train_1) #1 row
```

```
(array([[1.],
       [1.],
       [1.],
```

```

[1.],
[1.]]), 436)

X_train = np.hstack((X_train_1, X_train))
X_test = np.hstack((X_test_1, X_test))
X_train[:5]

array([[ 1.          ,  0.07687786,  0.0114873 ,  0.26301555,  1.          ,
         1.          , -0.21651786, -0.53954802, -0.04807692,  1.          ,
        0.13285601, -0.30695444],
       [ 1.          ,  0.18548046,  0.0114873 ,  0.26301555, -0.36705686,
         1.          , -0.21651786,  1.          , -0.04807692,  1.          ,
         1.          , -0.30695444],
       [ 1.          , -0.12077888, -0.31801693, -0.10547667, -0.36705686,
         1.          , -0.21651786,  1.          , -0.04807692,  1.          ,
        0.566428    , -0.30695444],
       [ 1.          , -0.2299245 ,  0.0114873 , -0.10547667,  0.08862876,
         1.          , -0.21651786,  1.          , -0.04807692, -0.46112601,
        -0.30071599,  1.          ],
       [ 1.          , -0.12675203,  0.0114873 , -0.10547667,  0.08862876,
         1.          , -0.21651786, -0.53954802, -0.04807692, -0.46112601,
        -0.30071599, -0.30695444]])

theta = np.zeros(12)
iterations = 1500
alpha = .01
theta.shape

(12,)

def compute_loss(X, Y, theta):
    predictions = X.dot(theta)
    errors = np.subtract(predictions, Y)
    sqrErrors = np.square(errors)
    J = 1/(2*m) * np.sum(sqrErrors)
    return J

cost = compute_loss(X_train,y_train, theta)
print('The cost for given values of theta_0 and theta_1 = ', cost)

The cost for given values of theta_0 and theta_1 =  0.07132905310217541

def gradient_descent(X, Y, theta, alpha, iterations,train = True):
    cost_history = np.zeros(iterations)

    for i in range(iterations):
        predictions = X.dot(theta)
        errors = np.subtract(predictions, Y)
        if train:
            sum_delta = (alpha/m)*X.transpose().dot(errors)
            theta = theta - sum_delta
            cost_history[i] = compute_loss(X, Y, theta)

    return theta, cost_history

theta, cost_history = gradient_descent(X_train, y_train,theta, alpha, iterations)
theta, test_history = gradient_descent(X_test, y_test,theta, alpha, iterations, train = False)
print('Final value of theta =', theta)
print('cost_history = ', cost_history[2:])
len(cost_history)

Final value of theta = [0.35585743 0.09012686 0.04931931 0.08952685 0.07344395 0.00611183
 0.02394733 0.02030264 0.03399489 0.04994878 0.06349352 0.03876027]
cost_history = [0.06731463 0.06604516 0.0648073 ... 0.00327253 0.0032722 0.00327188]
1500

plt.plot(cost_history, color = 'pink', label = 'training data')
plt.plot(test_history, color = 'blue', label = 'testing data')

```

[<matplotlib.lines.Line2D at 0x7f3770454130>]



Problem 3a

```

|-----|
#accessing columns that were asked
#loc = index by name and name is string
reduced_df = df.loc[:,["area", "bedrooms", "bathrooms", "stories", "parking"]]
price = df.iloc[:,0] #iloc = index by number

maxes = []
reduced_df = normalize_dataframe(reduced_df)
for i in range(5):
    maxes.append(max(reduced_df.iloc[:,i]))
    reduced_df.iloc[:,i] = reduced_df.iloc[:,i]/maxes[-1]

maxes.append(max(price))
price = price/maxes[-1]
reduced_df.iloc[:,0] = reduced_df.iloc[:,0]/max(reduced_df.iloc[:,0])
print(price)
print(reduced_df.iloc[:,:])
print(maxes)

0      1.000000
1      0.921053
2      0.921053
3      0.918421
4      0.857895
...
540    0.136842
541    0.132868
542    0.131579
543    0.131579
544    0.131579
Name: price, Length: 545, dtype: float64
   area  bedrooms  bathrooms  stories  parking
0  0.205391  0.340992   0.263016  0.544314  0.566428
1  0.344764  0.340992   1.000000  1.000000  1.000000
2  0.435266  0.011487   0.263016  0.088629  0.566428
3  0.212631  0.340992   0.263016  0.088629  1.000000
4  0.205391  0.340992  -0.105477  0.088629  0.566428
..      ...      ...      ...      ...      ...
540 -0.194629 -0.318017 -0.105477 -0.367057  0.566428
541 -0.248930  0.011487 -0.105477 -0.367057 -0.300716
542 -0.138517 -0.318017 -0.105477 -0.367057 -0.300716
543 -0.202774  0.011487 -0.105477 -0.367057 -0.300716
544 -0.117702  0.011487 -0.105477  0.088629 -0.300716

[545 rows x 5 columns]
[5.091585570242134, 4.111923842176972, 5.400846899372526, 2.529699688044935, 2.6769500506662025, 13300000]

X_train, X_test, y_train, y_test = train_test_split(reduced_df.values, price.values, test_size = 0.2, random_state =42)

print('features:\n', X_train[0]) #shows all 5 X values
print('example price:', y_train[0])

features:
[0.07687786 0.0114873  0.26301555 1.          0.13285601]
example price: 0.5657894736842105

# 5x5
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
m = len(X_train)
X_train = X_train.reshape(m,5)
print('Xvals: \n', X_train[:5,:])
m

Xvals:
[[ 0.07687786  0.0114873  0.26301555  1.          0.13285601]
```

```
[ 0.18548046  0.0114873  0.26301555 -0.36705686  1.          ]
[-0.12077888 -0.31801693 -0.10547667 -0.36705686  0.566428    ]
[-0.2299245   0.0114873  -0.10547667  0.08862876 -0.30071599]
[-0.12675203  0.0114873  -0.10547667  0.08862876 -0.30071599]]
436
```

```
m = len(X_train)
n = len(X_test)
X_train_1 = np.ones((m,1)) #1 column
X_test_1 = np.ones((n,1)) #1 column
X_train_1[:1], len(X_train_1) #1 row
```

```
(array([[1.]]), 436)
```

```
X_train = np.hstack((X_train_1, X_train))
X_test = np.hstack((X_test_1, X_test))
X_train[:5],
```

```
(array([[ 1.          ,  0.07687786,  0.0114873 ,  0.26301555,  1.          ,
          0.13285601],
[ 1.          ,  0.18548046,  0.0114873 ,  0.26301555, -0.36705686,
          1.          ],
[ 1.          , -0.12077888, -0.31801693, -0.10547667, -0.36705686,
          0.566428    ],
[ 1.          , -0.2299245 ,  0.0114873 , -0.10547667,  0.08862876,
         -0.30071599],
[ 1.          , -0.12675203,  0.0114873 , -0.10547667,  0.08862876,
         -0.30071599]]),)
```

```
theta = np.zeros(6)
iterations = 1500
alpha = .01
theta.shape
```

```
(6,)
```

```
def compute_loss(X, Y, theta, lambda_reg=0.1):
    m = len(Y)
    predictions = X.dot(theta)
    errors = predictions - Y
    sqrErrors = np.square(errors)
    regularization = (lambda_reg/(2*m)) * np.sum(theta[1:]**2)
    J = 1/(2*m) * np.sum(sqrErrors) + regularization
    return J
```

```
cost = compute_loss(X_train,y_train, theta)
print('The cost for given values of theta_0 and theta_1 = ', cost)
```

```
The cost for given values of theta_0 and theta_1 =  0.07132905310217541
```

```
def gradient_descent(X, Y, theta, alpha, iterations, lambda_reg=0.1, train=True):
    m = len(Y)
    cost_history = np.zeros(iterations)

    for i in range(iterations):
        predictions = X.dot(theta)
        errors = predictions - Y
        if train:
            regularization = (lambda_reg/(2*m)) * np.sum(theta[1:]**2)
            sum_delta = (alpha/m)*X.transpose().dot(errors) + regularization
            theta = theta - sum_delta
            cost_history[i] = compute_loss(X, Y, theta, lambda_reg)

    return theta, cost_history
```

```
X_train
```

```
array([[ 1.          ,  0.07687786,  0.0114873 ,  0.26301555,  1.          ,
          0.13285601],
[ 1.          ,  0.18548046,  0.0114873 ,  0.26301555, -0.36705686,
          1.          ],
[ 1.          , -0.12077888, -0.31801693, -0.10547667, -0.36705686,
          0.566428    ],
...,
[ 1.          , -0.0588754 ,  0.0114873 ,  0.26301555,  0.54431438,
```

```

0.13285601],
[ 1.          , -0.1005064 , -0.31801693, -0.10547667, -0.36705686,
 -0.30071599],
[ 1.          , 0.03162677, 0.0114873 , 0.26301555, 1.          ,
 0.13285601]])

```

X_train

```

array([[ 1.          , 0.07687786, 0.0114873 , 0.26301555, 1.          ,
        0.13285601],
 [ 1.          , 0.18548046, 0.0114873 , 0.26301555, -0.36705686,
        1.          ],
 [ 1.          , -0.12077888, -0.31801693, -0.10547667, -0.36705686,
        0.566428    ],
 ...,
 [ 1.          , -0.0588754 , 0.0114873 , 0.26301555, 0.54431438,
        0.13285601],
 [ 1.          , -0.1005064 , -0.31801693, -0.10547667, -0.36705686,
        -0.30071599],
 [ 1.          , 0.03162677, 0.0114873 , 0.26301555, 1.          ,
        0.13285601]])

```

```

theta, cost_history = gradient_descent(X=X_train, Y=y_train, theta=theta, alpha=alpha, iterations=iterations)
theta, test_history = gradient_descent(X_test, y_test, theta, alpha, iterations, train = False)

```

```

print('Final value of theta =', theta)
print('cost_history = ', cost_history[2:])
len(cost_history)

```

```

Final value of theta = [0.35545312 0.12310667 0.06147617 0.09879753 0.09257036 0.08708548]
cost_history = [0.06763439 0.0664511 0.06529114 ... 0.00445522 0.00445479 0.00445437]
1500


```

```

plt.plot(cost_history, color = 'pink', label = 'training data')
plt.plot(test_history, color = 'blue', label = 'testing data')

```

[<matplotlib.lines.Line2D at 0x7f3771bd7610>]



Problem 3b

```

#accessing columns that were asked
#loc = index by name and name is string
reduced_df = df.loc[:,["area", "bedrooms", "bathrooms", "stories", "mainroad", "guestroom", "basement", "hotwaterheating", "airconditioning",
price = df.iloc[:,0] #iloc = index by number

# arr = [1, 2, 3]
# print([i for i in arr])
# print([i**2 for i in arr])
# print([i for i in arr if i%2 == 1])

for column_index in [4, 5, 6, 7, 8, 10]:
    reduced_df.iloc[:,column_index] = reduced_df.iloc[:,column_index].map({"yes": 1, "no": 0})
display(reduced_df.head())

maxes = []
reduced_df = normalize_dataframe(reduced_df)
for i in range(11):
    maxes.append(max(reduced_df.iloc[:,i]))
    reduced_df.iloc[:,i] = reduced_df.iloc[:,i]/maxes[-1]

maxes.append(max(price))
price = price/maxes[-1]

```

```

reduced_df.iloc[:,0] = reduced_df.iloc[:,0]/max(reduced_df.iloc[:,0])
print(price)
print(reduced_df.iloc[:,:])
print(maxes)

```

	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating	airconditioning	parking	prefarea
0	7420	4	2	3	1	0	0	0	1	2	1
1	8960	4	4	4	1	0	0	0	1	3	0
2	9960	3	2	2	1	0	1	0	0	2	1
3	7500	4	2	2	1	0	1	0	1	3	1
4	7420	4	1	2	1	1	1	0	1	2	0



```

0    1.000000
1    0.921053
2    0.921053
3    0.918421
4    0.857895
...
540  0.136842
541  0.132868
542  0.131579
543  0.131579
544  0.131579
Name: price, Length: 545, dtype: float64
   area  bedrooms  bathrooms  stories  mainroad  guestroom  basement \
0  0.205391  0.340992  0.263016  0.544314  1.000000  -0.216518  -0.539548
1  0.344764  0.340992  1.000000  1.000000  1.000000  -0.216518  -0.539548
2  0.435266  0.011487  0.263016  0.088629  1.000000  -0.216518  1.000000
3  0.212631  0.340992  0.263016  0.088629  1.000000  -0.216518  1.000000
4  0.205391  0.340992  -0.105477  0.088629  1.000000  1.000000  1.000000
..  ...
540 -0.194629 -0.318017 -0.105477 -0.367057  1.000000  -0.216518  1.000000
541 -0.248930  0.011487 -0.105477 -0.367057 -0.077922  -0.216518  -0.539548
542 -0.138517 -0.318017 -0.105477 -0.367057  1.000000  -0.216518  -0.539548
543 -0.202774  0.011487 -0.105477 -0.367057 -0.077922  -0.216518  -0.539548
544 -0.117702  0.011487 -0.105477  0.088629  1.000000  -0.216518  -0.539548

   hotwaterheating  airconditioning  parking  prefarea
0      -0.048077      1.000000  0.566428  1.000000
1      -0.048077      1.000000  1.000000 -0.306954
2      -0.048077     -0.461126  0.566428  1.000000
3      -0.048077      1.000000  1.000000  1.000000
4      -0.048077      1.000000  0.566428 -0.306954
..  ...
540     -0.048077     -0.461126  0.566428 -0.306954
541     -0.048077     -0.461126 -0.300716 -0.306954
542     -0.048077     -0.461126 -0.300716 -0.306954
543     -0.048077     -0.461126 -0.300716 -0.306954
544     -0.048077     -0.461126 -0.300716 -0.306954

```

[545 rows x 11 columns]

```
X_train, X_test, y_train, y_test = train_test_split(reduced_df.values, price.values, test_size = 0.2, random_state =42)
```

```

print('features:\n', X_train[0]) #shows all 5 X values
print('example price:', y_train[0])

```

```

features:
[ 0.07687786  0.0114873  0.26301555  1.          1.          -0.21651786
 -0.53954802 -0.04807692  1.          0.13285601 -0.30695444]
example price: 0.5657894736842105

```

```

# 5x5
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
m = len(X_train)
X_train = X_train.reshape(m,11)
print('Xvals: \n', X_train[:11,:])
m

```

```

Xvals:
[[ 0.07687786  0.0114873  0.26301555  1.          1.          -0.21651786
 -0.53954802 -0.04807692  1.          0.13285601 -0.30695444]
 [ 0.18548046  0.0114873  0.26301555 -0.36705686  1.          -0.21651786
  1.          -0.04807692  1.          1.          -0.30695444]
 [-0.12077888 -0.31801693 -0.10547667 -0.36705686  1.          -0.21651786

```

```

1. -0.04807692 1. 0.566428 -0.30695444]
[-0.2299245 0.0114873 -0.10547667 0.08862876 1. -0.21651786
1. -0.04807692 -0.46112601 -0.30071599 1. ]
[-0.12675203 0.0114873 -0.10547667 0.08862876 1. -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.01271929 0.0114873 -0.10547667 0.08862876 1. -0.21651786
1. -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.11770181 0.0114873 -0.10547667 0.08862876 1. -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.00674615 0.0114873 -0.10547667 -0.36705686 -6.07792208 -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.17833826 0.0114873 -0.10547667 -0.36705686 -6.07792208 -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[ 0.0786879 0.0114873 -0.10547667 -0.36705686 1. -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.18376839 0.0114873 -0.10547667 0.08862876 -6.07792208 -0.21651786
1. 1. -0.46112601 -0.30071599 -0.30695444]]
436

```

```
# 5x5
```

```
# shows first 5 rows of area, bedrooms, bathrooms, stories, parking
```

```
m = len(X_train)
```

```
X_train = X_train.reshape(m,11)
```

```
print('Xvals: \n', X_train[:11,:])
```

```
m
```

```

Xvals:
[[ 0.07687786 0.0114873 0.26301555 1. 1. -0.21651786
-0.53954802 -0.04807692 1. 0.13285601 -0.30695444]
[ 0.18548046 0.0114873 0.26301555 -0.36705686 1. -0.21651786
1. -0.04807692 1. 1. -0.30695444]
[-0.12077888 -0.31801693 -0.10547667 -0.36705686 1. -0.21651786
1. -0.04807692 1. 0.566428 -0.30695444]
[-0.2299245 0.0114873 -0.10547667 0.08862876 1. -0.21651786
1. -0.04807692 -0.46112601 -0.30071599 1. ]
[-0.12675203 0.0114873 -0.10547667 0.08862876 1. -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.01271929 0.0114873 -0.10547667 0.08862876 1. -0.21651786
1. -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.11770181 0.0114873 -0.10547667 0.08862876 1. -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.00674615 0.0114873 -0.10547667 -0.36705686 -6.07792208 -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.17833826 0.0114873 -0.10547667 -0.36705686 -6.07792208 -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[ 0.0786879 0.0114873 -0.10547667 -0.36705686 1. -0.21651786
-0.53954802 -0.04807692 -0.46112601 -0.30071599 -0.30695444]
[-0.18376839 0.0114873 -0.10547667 0.08862876 -6.07792208 -0.21651786
1. 1. -0.46112601 -0.30071599 -0.30695444]]
436

```

```
m = len(X_train)
```

```
n = len(X_test)
```

```
X_train_1 = np.ones((m,1)) #1 column
```

```
X_test_1 = np.ones((n,1)) #1 column
```

```
X_train_1[:1], len(X_train_1) #1 row
```

```
(array([[1.]]), 436)
```

```
X_train = np.hstack((X_train_1, X_train))
```

```
X_test = np.hstack((X_test_1, X_test))
```

```
X_train[:5]
```

```

array([[ 1. , 0.07687786, 0.0114873, 0.26301555, 1. ,
1. , -0.21651786, -0.53954802, -0.04807692, 1. ,
0.13285601, -0.30695444],
[ 1. , 0.18548046, 0.0114873, 0.26301555, -0.36705686,
1. , -0.21651786, 1. , -0.04807692, 1. ,
1. , -0.30695444],
[ 1. , -0.12077888, -0.31801693, -0.10547667, -0.36705686,
1. , -0.21651786, 1. , -0.04807692, 1. ,
0.566428, -0.30695444],
[ 1. , -0.2299245, 0.0114873, -0.10547667, 0.08862876,
1. , -0.21651786, 1. , -0.04807692, -0.46112601,
-0.30071599, 1. ],
[ 1. , -0.12675203, 0.0114873, -0.10547667, 0.08862876,
1. , -0.21651786, -0.53954802, -0.04807692, -0.46112601,
-0.30071599, -0.30695444]])

```

```
theta = np.zeros(12)
iterations = 1500
alpha = .01
theta.shape
```

```
(12,)
```

```
def compute_loss(X, Y, theta):
    predictions = X.dot(theta)
    errors = np.subtract(predictions, Y)
    sqrErrors = np.square(errors)
    J = 1/(2*m) * np.sum(sqrErrors)
    return J
```

```
cost = compute_loss(X_train,y_train, theta)
print('The cost for given values of theta_0 and theta_1 = ', cost)
```

```
The cost for given values of theta_0 and theta_1 = 0.07132905310217541
```

```
def gradient_descent(X, Y, theta, alpha, iterations):
    cost_history = np.zeros(iterations)
```

```
    for i in range(iterations):
        predictions = X.dot(theta)
        errors = np.subtract(predictions, Y)
        sum_delta = (alpha/m)*X.transpose().dot(errors)
        theta = theta - sum_delta
        cost_history[i] = compute_loss(X, Y, theta)
```


```
    return theta, cost_history
```

```
theta, cost_history = gradient_descent(X_train, y_train,theta, alpha, iterations)
print('Final value of theta =', theta)
print('cost_history = ', cost_history[2:])
len(cost_history)
```

```
Final value of theta = [0.35585743 0.09012686 0.04931931 0.08952685 0.07344395 0.00611183
 0.02394733 0.02030264 0.03399489 0.04994878 0.06349352 0.03876027]
cost_history = [0.06731463 0.06604516 0.0648073 ... 0.00327253 0.0032722 0.00327188]
1500
```

```
plt.plot(cost_history, color = 'pink', label = 'training data')
plt.plot(test_history, color = 'blue', label = 'testing data')
```

```
[<matplotlib.lines.Line2D at 0x7f37706af5e0>]
```



✓ 0s completed at 11:40 PM

● ✕